

Data Containers

Francesc Alted

Freelance Consultant

<http://www.blosc.org/professional-services.html>

Advanced Scientific Programming in Python

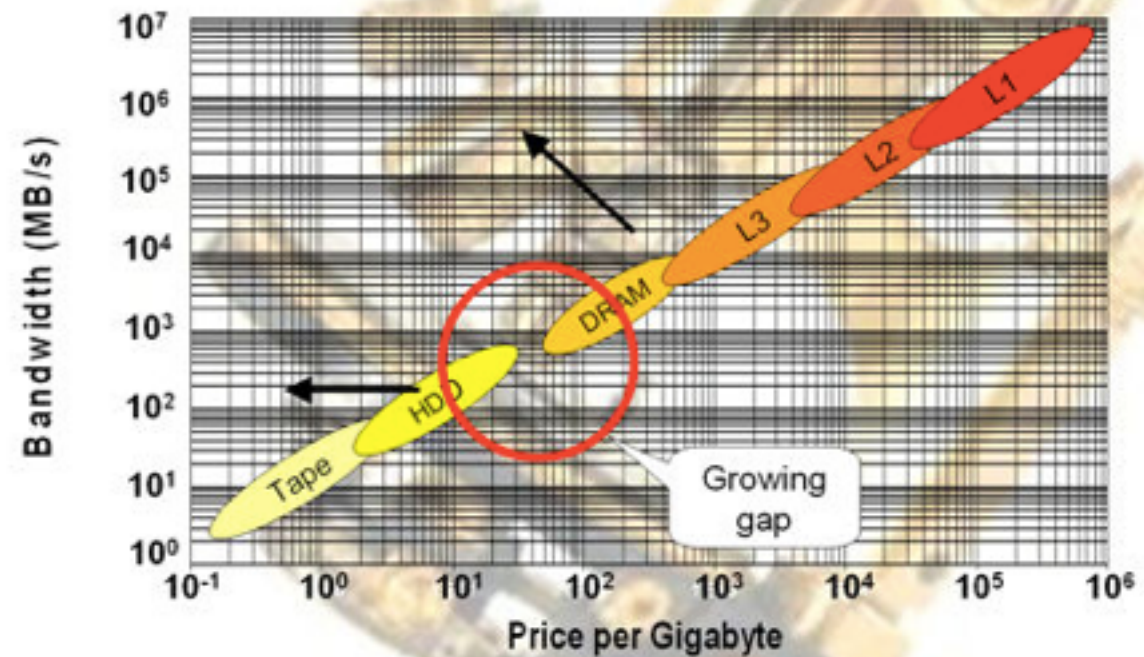
Reading, UK

September, 2016


Trends in Computer Storage

The growing gap between DRAM and HDD is facilitating the introduction of new SDD devices

The DRAM/HDD Speed Gap



From: *Solid State Drives in the Enterprise*
by Objective Analysis



PCIe SSD



M.2 SSD



BGA SSD

Latency Numbers Every Programmer Should Know


Latency Comparison Numbers

L1 cache reference	0.5	ns			
Branch mispredict	5	ns			
L2 cache reference	7	ns			14x L1 cache
Mutex lock/unlock	25	ns			
Main memory reference	100	ns			20x L2 cache, 200x L1 cache
Read 4K randomly from memory	1,000	ns	0.001	ms	
Compress 1K bytes with Zippy	3,000	ns			
Send 1K bytes over 1 Gbps network	10,000	ns	0.01	ms	
Read 4K randomly from SSD*	150,000	ns	0.15	ms	
Read 1 MB sequentially from memory	250,000	ns	0.25	ms	
Round trip within same datacenter	500,000	ns	0.5	ms	
Read 1 MB sequentially from SSD*	1,000,000	ns	1	ms	4X memory
Disk seek	10,000,000	ns	10	ms	20x datacenter roundtrip
Read 1 MB sequentially from disk	20,000,000	ns	20	ms	80x memory, 20X SSD
Send packet CA->Netherlands->CA	150,000,000	ns	150	ms	

Source: Jeff Dean and Peter Norvig (Google), with some additions

http://www.eecs.berkeley.edu/~rcs/research/interactive_latency.html

Reference Time vs Transmission Time



$t_{ref} \sim t_{trans} \Rightarrow$ optimizes memory access

Not All Storage Layers Are Created Equal

Memory: t_{ref} : 100 ns / t_{trans} (**1 KB**): ~100 ns

Solid State Disk: t_{ref} : 10 μ s / t_{trans} (**4 KB**): ~10 μ s

Mechanical Disk: t_{ref} : 10 ms / t_{trans} (**1 MB**): ~10 ms

The slower the media, the larger the block that is worth to transmit

But essentially, a blocked data access is mandatory for speed!

We Need More Data Blocking In Our Infrastructure!

- Not many data containers leveraging the blocking technique exist yet, but a handful (e.g. HDF5, bcolz or zarr) do
- With blocked access we can use persistent media (disk) as it is ephemeral (memory) and the other way around -> independency of media!
- **No silver bullet:** we won't be able to find a single container that makes everybody happy; it's all about tradeoffs

Can We Get Better
Bandwidth Than
Hardware Allows?

Compression for Random & Sequential Access in SSDs

Performance Specification	Incompressible Data	Compressible Data
Sequential Write Bandwidth (Mbps)	235	520
Sequential Read Bandwidth (Mbps)	550	550
Random Write (IOPS)	16,500 (65MB/s)	60,000 (240MB/s)
Random Read (IOPS)	46,000 (180MB/s)	50,000 (200MB/s)

3. Source: Intel® Solid-State Drive 520 Series Product Specification; Random reads based on 4KB Queue Depth 32

- Compression does help performance!

Compression for Random & Sequential Access in SSDs


Performance Specification	Incompressible Data	Compressible Data
Sequential Write Bandwidth (Mbps)	235	520
Sequential Read Bandwidth (Mbps)	550	550
Random Write (IOPS)	16,500 (65MB/s)	60,000 (240MB/s)
Random Read (IOPS)	46,000 (180MB/s)	50,000 (200MB/s)

3. Source: Intel® Solid-State Drive 520 Series Product Specification; Random reads based on 4KB Queue Depth 32


- Compression does help performance!
- However, limited by SATA bandwidth

Leveraging Compression Straight To CPU

Less data needs to be transmitted to the CPU




Transmission + decompression faster than direct transfer?



When we have a fast enough compressor we can get rid of the limitations of the bus bandwidth.


How to get maximum compression performance?

NETCDF C LIB: Write TP vs File Size cksize_T:23 cksize_YX:128



Thanks to: Rui Yang, Pablo Larraondo (NCI Australia)

Example with actual data (satellite images):
Blosc compression does not degrade I/O performance




Thanks to: Rui Yang, Pablo Larraondo (NCI Australia)

Reading satellite images:
Blosc decompression accelerates I/O

Can CPU-based
Compression Alleviate
The Memory Bottleneck?

Improving RAM Speed?

Less data needs to be transmitted to the CPU



Transmission + decompression faster than direct transfer?


We can try, but certain conditions must be met

- We must follow the principles we have seen:
 - Data Containers leveraging the blocking technique (better cache usage)
 - Using (extremely fast) compression per every block

Principles of Blosc


- Split data chunks in blocks internally (better cache utilization)
- Supports Shuffle and BitShuffle filters
- Uses parallelism at two levels:
 - Use multicores (multithreading)
 - Use SIMD in Intel/AMD processors (SSE2, AVX2) and ARM (NEON)

The Shuffle filter



- Shuffle works at byte level, and works well for integers or floats that vary smoothly
- There is also support for a BitShuffle filter that works at bit level

Blosc: Compressing Faster Than *memcpy()*




bcolz: a Data Container that Leverages the Blocking Technique

NumPy container



Contiguous memory

bcolz container

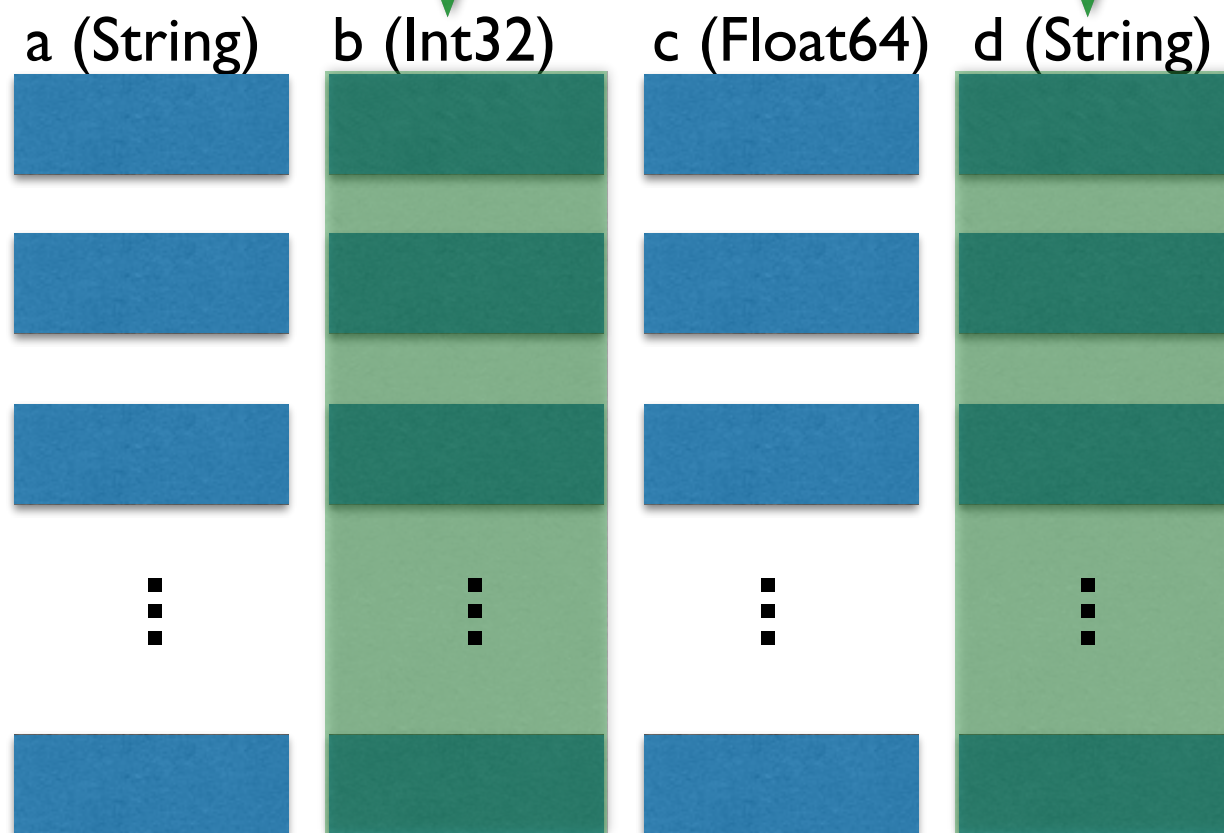


Discontiguous memory

Interesting columns




Interesting rows

Chunk 1



Chunked Query

Chunk N




Query: `(b == 5) & (d == 'some string')`

Very efficient when query
selectivity is high and
decompression is fast

Query Times in bcolz

Recent server (Intel Xeon Skylake, 4 cores)
Compression **speeds** things up




Reference: <https://github.com/Blosc/movielens-bench/blob/master/querying-ep14.ipynb>

Query Times in bcolz


4-year old laptop (Intel Ivy-Bridge, 2 cores)

Compression still **speeds** things up




Query Times in bcolz

2010 laptop (Intel Core2, 2 cores)
Compression **slows** things down




Sizes in bcolz



Do not forget compression main strength:
We can store more data while using same resources

Accelerating I/O With Blosc



“Blosc compressors are the fastest ones out there at this point; there is no better publicly available option that I'm aware of. That's not just ‘yet another compressor library’ case.”

— *Ivan Smirnov*
(*advocating for Blosc inclusion in h5py*)

Compression matters!

Closing Notes

- Due to the evolution in computer architecture, compression can be effective for two reasons:
 - We can work with more data using the same resources.
 - We can reduce the overhead of compression to near zero, and even beyond than that!

Thanks!